

FRAMEBUFFERS

The combination of the color buffer, depth buffer, and stencil buffer is stored in GPU memory known as the **framebuffer**. OpenGL gives us the flexibility to define our own **framebuffers** and, thus, define our own color (and optionally a depth and stencil) buffer.

The rendering options we've done so far were all done on top of the render buffers attached to the **default framebuffer**.

- The **default framebuffer** is created and configured when you create your window (GLFW does this for us). By creating our own **framebuffer**, we can get an additional target to render to.
- Rendering your scene to a different **framebuffer** allows us to use that result create mirrors in a scene, add post-processing effects, etc.

Creating a Framebuffer

Generate and bind an OpenGL **framebuffer** object.

- [Generating and Binding a Framebuffer](#)

NOTE: It is possible to bind a **framebuffer** to a read or write target specifically by binding to `GL_READ_FRAMEBUFFER` or `GL_DRAW_FRAMEBUFFER`, respectively.

- The **framebuffer** bound to `GL_READ_FRAMEBUFFER` is used for all read operations like `glReadPixels`.
- The **framebuffer** bound to `GL_DRAW_FRAMEBUFFER` is used as the destination for rendering, clearing, and other write operations.

We cannot use our **framebuffer** yet because it is not **complete**. To complete the **framebuffer**, the following requirements must be satisfied:

1. At least one buffer is attached.
2. There is at least one color attachment.
3. All attachments are complete as well (reserved memory).
4. Each buffer has the same number of samples (samples are discussed in a later chapter).

We should use `glCheckFramebufferStatus(GL_FRAMEBUFFER)` to check if the **framebuffer** is complete.

- `glCheckFramebufferStatus` checks the **framebuffer** for completeness and returns the completeness status of the **framebuffer**.
- [Specification for glCheckFramebufferStatus](#)

Since our **framebuffer** is not the **default framebuffer**, the rendering commands will have no impact on the visual output of your window.

- For this reason, it is called **off-screen rendering** when rendering to a different **framebuffer**.
- In order to make all rendering options have a visual impact again on the main window, you must re-bind the **default framebuffer** (0).

Check if the **framebuffer** is complete. Then, re-bind the **default framebuffer**.

- Don't forget to delete the **framebuffer** object you generated after you are done using it.
- [Checking Framebuffer Completeness](#)

The first requirement for **framebuffer** completeness states that we need to attach one or more **attachments** to the **framebuffer**.

- An **attachment** is a memory location that can act as a buffer for the **framebuffer** (think of it as an image).
- When creating an **attachment**, we have two options:
 - **Textures**
 - **Renderbuffer objects**

Texture Attachments

When attaching a **texture** to a **framebuffer**, all rendering commands will write to the **texture** as if it was a normal color/depth/stencil buffer.

- The advantage of using **textures** is that the render output is stored inside the **texture image**, which we can easily use in our shaders.

Creating a **texture** for a **framebuffer** is the same as creating a normal **texture**, except you don't (typically) specify wrapping or mipmaps, and you use the window size as the image size parameters.

You can attach a **2D texture** to a **framebuffer** using

`glFramebufferTexture2D(GLenum target, GLenum textarget, GLuint texture, GLint level)`.

- **target:** The **framebuffer** type we're targeting (draw, read, or both)
- **attachment:** The type of **attachment** we're going to attach. Right now, we're attaching a **color attachment**.
- **textarget:** The type of the **texture** you want to attach.
- **texture:** The **texture** to attach.
- **level:** The mipmap level.

Create an empty **texture** object with linear filtering and attach it to the **framebuffer**.

- [Attaching a Texture](#)

Next to the **color attachment**, we can also attach a **depth texture** (using `GL_DEPTH_ATTACHMENT`), which would cause the **texture's** format and internalformat types to become `GL_DEPTH_COMPONENT`, and a **stencil texture** (using `GL_STENCIL_ATTACHMENT`), which would cause the formats to become `GL_STENCIL_INDEX`.

It is also possible to attach both a depth buffer and a stencil buffer as a single **texture**.

- Each 32-bit value of the **texture** then contains 24 bits of depth information and 8 bits of stencil information.
- To attach a depth and stencil buffer as one **texture**, we use the `GL_DEPTH_STENCIL_ATTACHMENT` type and configure the **texture's** formats to contain combined depth and stencil values.
 - [Example of Attaching a Depth-Stencil Texture](#)

Renderbuffer Object Attachments

Just like a **texture image**, a **renderbuffer object** is an actual buffer (e.g. an array of bytes, or integers, or pixels, etc.). However, a **renderbuffer object** cannot be directly read from.

- Not being able to be directly read from gives it the added advantage that OpenGL can do a few memory optimizations that can give it a performance edge over **textures** for off-screen rendering to a **framebuffer**.

Renderbuffer objects store all the render data directly into their buffer without any conversions to **texture**-specific formats, making them faster as a writeable storage medium.

- You can't read from a **renderbuffer object** directly, but you can use `glReadPixels` (very slow) to return a specified area of pixels from the currently bound **framebuffer**. This, however, does not return data directly from the **attachment** itself.

Because their data is in a native format, they are quite fast when writing data or copying data to other buffers.

- Operations like switching buffers are, therefore, quite fast when using **renderbuffer objects**.
- The `glFwSwapBuffers` function we've been using at the end of each frame may be implemented with **renderbuffer objects** (we simply write to a **renderbuffer image** and swap to the other one at the end). **Renderbuffer objects** are perfect for these kinds of operations.

Since **renderbuffer objects** are write-only, they are often used as **depth and stencil attachments**, since, most of the time, we don't really need to read values from them, but we care about depth and stencil testing.

- We need the depth and stencil values for testing, but we don't need to sample these values. So, a **renderbuffer object** suits this perfectly.
- When we're not sampling from these buffers, a **renderbuffer object** is generally preferred.

Creating a **depth and stencil renderbuffer object** is done by calling the `glRenderbufferStorage` function. You can then attach the **renderbuffer object** using the `glFramebufferRenderbuffer` function.

Create a **depth and stencil renderbuffer object** and attach it to the **framebuffer**.

- [Attaching a Renderbuffer Object](#)

It is important to realize when to use **renderbuffer objects** and when to use **textures**.

- The general rule is that:
 - If you never need to sample data from a specific buffer, it is wise to use a **renderbuffer object** for that specific buffer.
 - However, if you need to sample data from a specific buffer (like color or depth values), you should use a **texture attachment** instead.

Rendering to a Texture

We're going to render the scene into a **texture** attached to a **framebuffer** object, and then draw this **texture** over a simple quad that spans the whole screen.

With a **texture color buffer** and a **depth-stencil renderbuffer object** attached, our **framebuffer** object is now complete.

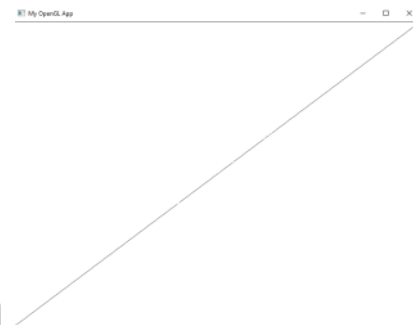

- **NOTE:** The **framebuffer** object was complete as soon as we performed the **texture color buffer attachment**, but we wanted an **attachment** for handling depth and stencil testing, so we attached a **depth-stencil renderbuffer object** as well.
- Now, all we need to do to render to our **framebuffer's** buffers instead of the **default framebuffer's** is to bind the **framebuffer** object.
 - **NOTE:** If you were to omit a **depth buffer attachment** for your **framebuffer**, depth testing operations would no longer work.

To draw the scene to a single **texture**:

1. Render the scene as usual with the new **framebuffer** bound as the active **framebuffer**.
2. Bind to the **default framebuffer**.
3. Draw a quad that spans the entire scene with the new **framebuffer's** color buffer as its **texture**.

Execute the steps above to draw the scene to a single **texture** that is drawn on a quad that covers to whole screen.

- You can test whether you did this successfully by changing the drawing mode to `GL_LINE` with `glPolygonMode(GL_FRONT_AND_BACK, GL_LINE)`. If the whole scene is just two triangles, but looks normal when not drawn as a wireframe, then you've implemented this successfully.
- [Rendering to a Texture](#)



Post-Processing

Now that the entire scene is rendered to a single **texture**, we can create post-processing effects by manipulating the scene **texture**.


For the following sections, I recommend you apply the container texture from previous chapters (or any texture, really) to the cube in the scene so you can better see the effects of post-processing.

Inversion

We have access to each of the colors of the render output, so it's not difficult to return the inverse of these colors in the fragment shader. You can accomplish this by taking the color of the screen **texture** and inverting it by subtracting it from `1.0`.

Modify your screen shader so that the colors displayed are inverted.

- [Inverted Colors](#)




Grayscale

Another interesting effect is to remove all colors from the scene except white, gray, and black color, converting the displayed image to grayscale. This involves setting the RGB components of the fragment to the average of those components.

- **NOTE:** The human eye tends to be more sensitive to green colors, and less sensitive to blue. To get more realistic results, adjust the color values before calculating the average of the colors to reflect how the human eye receives color.

Modify your screen shader so that the colors displayed are in grayscale.

- [Grayscale Colors](#)



Kernel Effects

Another advantage of doing post-processing on a single **texture image** is that we can sample color values from other parts of the **texture** not specific to that fragment.

- **Example:** We could take a small area around the current **texture** coordinate and sample multiple **texture** values around the current **texture** value. We could then create interesting effects by combining them in different ways.

A **kernel** (or convolution matrix) is a small, matrix-like array of values centered on the current pixel that multiplies the surrounding pixel values by its **kernel** values and adds them all together to form a single value. We add a small offset to the texture coordinates in surrounding directions of the current pixel and combine the results based on the **kernel**.

interesting effects by combining them in different ways.

A **kernel** (or convolution matrix) is a small, matrix-like array of values centered on the current pixel that multiplies the surrounding pixel values by its **kernel** values and adds them all together to form a single value. We add a small offset to the texture coordinates in surrounding directions of the current pixel and combine the results based on the **kernel**.

The sample **kernel** in the image to the right takes 8 surrounding pixel values and multiplies them by 2 and the current pixel by -15 and sharpens the image.

$$\begin{bmatrix} 2 & 2 & 2 \\ 2 & -15 & 2 \\ 2 & 2 & 2 \end{bmatrix}$$

- This **kernel** multiplies the surrounding pixels by several weights determined in the **kernel** and balances the result by multiplying the current pixel by a large negative weight.
- **NOTE:** Most **kernels** you'll find over the internet all sum up to 1 if you add all the weights together. If they don't add up to 1, it means that the resulting **texture** color ends up brighter or darker than the original **texture** value.

Kernels are an extremely useful tool for post-processing since they're easy to experiment with and have a lot of examples available online.

Modify your screen shader to implement a *sharpen kernel*, which sharpens the image.

- $\begin{bmatrix} -1 & -1 & -1 \\ -1 & 9 & -1 \\ -1 & -1 & -1 \end{bmatrix}$ is an example of a sharpen kernel.
- [Sharpen Kernel](#)

All you have to change is the **kernel**, and you can get very different post-processing effects. This makes using **kernels** very easy and powerful.

Blur




Modify your screen shader to implement a *blur kernel*, which blurs the image.

- $\begin{bmatrix} 1 & 2 & 1 \\ 2 & 4 & 2 \\ 1 & 2 & 1 \end{bmatrix} / 16$ is an example of a blur kernel.
- [Blur Kernel](#)

Edge Detection

Modify your screen shader to implement an *edge-detection kernel*, which highlights edges (color-wise) and darkens the rest.

- $\begin{bmatrix} 1 & 1 & 1 \\ 1 & 8 & 1 \\ 1 & 1 & 1 \end{bmatrix}$ is an example of an edge-detection kernel.
- [Edge-Detection Kernel](#)



EXERCISES

1. Can you use framebuffers to create a rear-view mirror? For this, you'll have to draw your scene twice: once with the camera rotated 180 degrees and once as normal. Try to create a small quad at the top of your screen to apply the mirror texture on.
2. Play around with the kernel values and create your own interesting post-processing effects. Try searching the internet as well for other interesting kernels.

Generating and Binding a Framebuffer

Friday, April 29, 2022 2:49 PM

```
unsigned int fbo;  
glGenFramebuffers(1, &fbo);  
glBindFramebuffer(GL_FRAMEBUFFER, fbo);
```

Checking Framebuffer Completeness

Friday, April 29, 2022 4:18 PM

```
if (glCheckFramebufferStatus(GL_FRAMEBUFFER) == GL_FRAMEBUFFER_COMPLETE) {  
    // framebuffer is complete  
}  
  
glBindFramebuffer(GL_FRAMEBUFFER, 0);  
  
...  
  
glDeleteFramebuffers(1, &fbo);
```

Attaching a Texture

Friday, April 29, 2022 5:26 PM

```
// Generates and binds a texture object
unsigned int texture;
glGenTextures(1, &texture);
glBindTexture(GL_TEXTURE_2D, texture);

// Image size set to size of screen. Data set to null since we're just allocating memory, not filling it.
glTexImage2D(GL_TEXTURE_2D, 0, GL_RGB, WIDTH, HEIGHT, 0, GL_RGB, GL_UNSIGNED_BYTE, NULL);

glTexParameteri(GL_TEXTURE_2D, GL_TEXTURE_MIN_FILTER, GL_LINEAR);
glTexParameteri(GL_TEXTURE_2D, GL_TEXTURE_MAG_FILTER, GL_LINEAR);

// We're done configuring this texture object, so we can unbind it.
glBindTexture(GL_TEXTURE_2D, 0);

// Attaches the texture to the framebuffer as a color attachment
glFramebufferTexture2D(GL_FRAMEBUFFER, GL_COLOR_ATTACHMENT0, GL_TEXTURE_2D, texture, 0);
```

NOTE: The 0 at the end of GL_COLOR_ATTACHMENT0 suggests we can attach more than 1 color attachment.

Example of Attaching a Depth-Stencil Texture

Friday, April 29, 2022 5:47 PM

```
// Specifies a texture composed of 24 bits of depth info and 8 bits of stencil info.  
glTexImage2D(GL_TEXTURE_2D, 0, GL_DEPTH24_STENCIL8, 800, 600, 0, GL_DEPTH_STENCIL, GL_UNSIGNED_INT_24_8, NULL);  
  
// Attaches the depth-stencil texture to the framebuffer  
glFramebufferTexture2D(GL_FRAMEBUFFER, GL_DEPTH_STENCIL_ATTACHMENT, GL_TEXTURE_2D, texture, 0);
```

Attaching a Renderbuffer Object

Friday, April 29, 2022 6:32 PM

```
// Generates and binds a renderbuffer object
unsigned int rbo;
glGenRenderbuffers(1, &rbo);
glBindRenderbuffer(GL_RENDERBUFFER, rbo);

// Makes the RBO a depth and stencil RBO, with 24 bits allocated for depth info, and 8 bits for stencil info
glRenderbufferStorage(GL_RENDERBUFFER, GL_DEPTH24_STENCIL8, WIDTH, HEIGHT);

// We're done configuring this renderbuffer object, so we can unbind it.
glBindRenderbuffer(GL_RENDERBUFFER, 0);

// Attaches the RBO to the framebuffer as a depth and stencil attachment
glFramebufferRenderbuffer(GL_FRAMEBUFFER, GL_DEPTH_STENCIL_ATTACHMENT, GL_RENDERBUFFER, rbo);
```


Rendering to a Texture

Friday, April 29, 2022 8:28 PM

screen.vert

```
#version 330 core

// aPos is a vec2 because the vertex position will be in normalized device coordinates (screen space)
layout (location = 0) in vec2 aPos;
layout (location = 1) in vec2 aTexCoords;

out vec2 texCoords;

void main() {
    gl_Position = vec4(aPos.x, aPos.y, 0.0, 1.0);
    texCoords = aTexCoords;
}
```

screen.frag

```
#version 330 core

in vec2 texCoords;

out vec4 FragColor;

uniform sampler2D screenTexture;

void main() {
    FragColor = texture(screenTexture, texCoords);
}
```

main.cpp

```
float screenQuadVertices[] = {
    // Positions      // Tex Coords
    -1.0f, -1.0f,      0.0f, 0.0f,      // bottom left
    -1.0f,  1.0f,      0.0f, 1.0f,      // top left
     1.0f,  1.0f,      1.0f, 1.0f,      // top right

    -1.0f, -1.0f,      0.0f, 0.0f,      // bottom left
     1.0f, -1.0f,      1.0f, 0.0f,      // bottom right
     1.0f,  1.0f,      1.0f, 1.0f,      // top right
};
```

main()

```
Shader screenShader{ "Shaders/screen.vert", "Shaders/screen.frag" };
...
// Screen quad VAO setup
unsigned int screenQuadVAO, screenQuadVBO;
glGenVertexArrays(1, &screenQuadVAO);
glGenBuffers(1, &screenQuadVBO);
glBindVertexArray(screenQuadVAO);
glBindBuffer(GL_ARRAY_BUFFER, screenQuadVBO);
glBufferData(GL_ARRAY_BUFFER, sizeof(screenQuadVertices), screenQuadVertices, GL_STATIC_DRAW);
glVertexAttribPointer(0, 2, GL_FLOAT, GL_FALSE, 4 * sizeof(float), (void*)0);
glVertexAttribPointer(1, 2, GL_FLOAT, GL_FALSE, 4 * sizeof(float), (void*)(sizeof(float) * 2));
glEnableVertexAttribArray(0);
glEnableVertexAttribArray(1);
glBindVertexArray(0);
glBindBuffer(GL_ARRAY_BUFFER, 0);

// FBO setup (in previous pages)
[...]
```

Render Loop

```
// RENDERING
// Bind our framebuffer object to be drawn to, with color and depth testing
glBindFramebuffer(GL_FRAMEBUFFER, fbo);
glClearColor(0.2f, 0.2f, 0.2f, 1.0f);
glClear(GL_COLOR_BUFFER_BIT | GL_DEPTH_BUFFER_BIT);
// Enables depth testing, since it is disabled later in a previous loop
glEnable(GL_DEPTH_TEST);
// Draw Scene
[...]
```



```
// Bind default framebuffer with white background color, with just color
glBindFramebuffer(GL_FRAMEBUFFER, 0);
```

```
glClearColor(1.0f, 1.0f, 1.0f, 1.0f);
glClear(GL_COLOR_BUFFER_BIT);

// Draws the color buffer texture to the screen, in front of everything else
screenShader.use();
glBindVertexArray(screenQuadVAO);
glDisable(GL_DEPTH_TEST);
glBindTexture(GL_TEXTURE_2D, textureColorBuffer);
glDrawArrays(GL_TRIANGLES, 0, 6);

glfwSwapBuffers(window);
glfwPollEvents();
```

NOTE: Since each framebuffer we're using has its own set of buffers, we want to clear each of those buffers with the appropriate bits set by calling `glClear`. Also, when drawing the quad, we're disabling depth testing since we want to make sure the quad always renders in front of everything else (we'll have to enable depth testing again when we draw the normal scene though).

Inverted Colors

Monday, May 2, 2022 1:26 PM

screen.frag

```
void main() {  
    FragColor = vec4(vec3(1.0 - texture(screenTexture, texCoords)), 1.0);  
}
```

We don't want to modify the w component when implementing this effect, so we cast `1.0 - texture(screenTexture, texCoords)` to a `vec3`, then cast it back to a `vec4` with `1.0` as the w component.

Grayscale Colors

Monday, May 2, 2022 1:35 PM

screen.frag

```
void main() {  
    FragColor = texture(screenTexture, texCoords);  
    float average = (FragColor.r + FragColor.g + FragColor.b) / 3.0; // True grayscale  
    FragColor = vec4(average, average, average, 1.0);  
}
```


To get a more realistic-looking grayscale, use the color calculations below.

- **NOTE:** The difference between the two might not be immediately visible, but more complex scene will look more realistic with the grayscale calculation used below as opposed to true grayscale.


screen.frag

```
void main() {  
    FragColor = texture(screenTexture, texCoords);  
    float average = FragColor.r * 0.2126 + FragColor.g * 0.7152 + FragColor.b * 0.0722; // More realistic grayscale  
    FragColor = vec4(average, average, average, 1.0);  
}
```

True Grayscale



More Realistic Grayscale



Sharpen Kernel

Monday, May 2, 2022 3:53 PM

```
...

const float offset = 1.0 / 300.0;

void main() {
    // Offset values that are used in determining what adjacent fragments to sample
    vec2 offsets[9] = vec2[(
        vec2(-offset, offset), // top-left
        vec2( 0.0f,   offset), // top-center
        vec2( offset, offset), // top-right
        vec2(-offset, 0.0f),   // center-left
        vec2( 0.0f,   0.0f),   // center-center
        vec2( offset, 0.0f),   // center-right
        vec2(-offset, -offset), // bottom-left
        vec2( 0.0f,   -offset), // bottom-center
        vec2( offset, -offset)  // bottom-right
    )];

    // Sharpen kernel
    float kernel[9] = float[(
        -1, -1, -1,
        -1, 9, -1,
        -1, -1, -1
    )];

    // Stores the adjacent fragment colors for kernelling
    vec3 sampleTex[9];
    for (int i = 0; i < 9; ++i) {
        sampleTex[i] = vec3(texture(screenTexture, texCoords.st + offsets[i]));
    }

    // Applies the kernel to the adjacent fragments' colors to determine this fragment's color
    vec3 color = vec3(0.0);
    for (int i = 0; i < 9; ++i) {
        color += sampleTex[i] * kernel[i];
    }

    FragColor = vec4(color, 1.0);
}
```

You can modify the offset value to influence the accuracy of the sharpen.

Blur Kernel

Monday, May 2, 2022 4:11 PM

```
// Blur kernel
float kernel[9] = float[(
    1.0 / 16.0, 2.0 / 16.0, 1.0 / 16.0,
    2.0 / 16.0, 4.0 / 16.0, 2.0 / 16.0,
    1.0 / 16.0, 2.0 / 16.0, 1.0 / 16.0
)];
```

$$\begin{bmatrix} 1 & 2 & 1 \\ 2 & 4 & 2 \\ 1 & 2 & 1 \end{bmatrix}$$

Because all values add up to 16, directly returning the combined sampled colors would result in an extremely bright color, so we have to divide each value of the kernel by 16.

Edge-Detection Kernel

Monday, May 2, 2022

4:25 PM

```
// Edge-detection kernel
float kernel[9] = float[](
    1.0,  1.0, 1.0,
    1.0, -8.0, 1.0,
    1.0,  1.0, 1.0
);
```